

Agenda

11:45 - 12:30: Brief introduction to Parallel Computing with OpenMP - Session 2

- OpenMP Work sharing (45 min + Hands On)

OpenMP Work Sharing

- Parallel region creates an *Single Program Multiple Data* instance where each thread executes the same code
- we can one split the work between the threads of a parallel region?
 - Loop construct
 - Task construct

OpenMP Work Sharing

- Directive instructing compiler to share the work of a loop

- C/C++: `#pragma omp for [clauses]`

- Fortran: `!$omp do [clauses]`

- The construct must be followed by a loop construct. To be active it must be inside a parallel region
- Combined construct with parallel:

in C/C++:

```
#pragma omp parallel for
```

in Fortran

```
$omp parallel do
```

- Loop index is private by default
- Work sharing dynamics can be controlled with the `schedule` clause

For-loops

- in C/C++

```
#pragma omp parallel for  
for(int i=0; i<n; i++){  
    ...  
}
```

- in Fortran

```
!$OMP PARALLEL DO  
do i = 1, n  
    ...  
end do  
!$OMP END PARALLEL DO
```

OpenMP Work Sharing

Loop construct

```
#pragma omp parallel for
  for(int i=0; i<n; i++){
    ...
  }
```

the n iterations will be split over the available threads accordingly

- Static scheduling, e.g.:

```
#pragma omp parallel for schedule(static, 10)
```

a chunk is 10 iterations. Threads receive a chunk to work in order.

- Dynamic scheduling, e.g.:

```
#pragma omp parallel for schedule(dynamic, 10)
```

a chunk is 10 iterations. Threads receive a chunk to work until they are exhausted.

- Guided scheduling, e.g.:

```
#pragma omp parallel for schedule(guided)
```

chunk size is modified as iterations are consumed.

OpenMP Work Sharing

Loop construct

- For loops

```
#pragma omp parallel
{
  #pragma omp for
  for(int i=0; i<n; i++){
    ...
  }
}
```

use within a parallel region, i.e. when a parallel region is already open.

- Race condition:

```
int sum_variable = 0;
#pragma omp parallel for
for(int i=0; i<n; i++){
  sum_variable += ...;
  ...
}
```

- takes place when multiple threads read and write a variable simultaneously
- random results depending on the order of threads accessing `sum_variable`

OpenMP Work Sharing

Reductions

- Summing elements of array is an example of reduction operation

$$S = \sum_{j=1}^N A_j = B_1 + B_2$$

- OpenMP provides support for common reductions within parallel regions and loops with the `reduction` -clause

OpenMP Work Sharing

Reductions

- Summing elements of array is an example of reduction operation

$$S = \sum_{j=1}^N A_j = B_1 + B_2$$

- OpenMP provides support for common reductions within parallel regions and loops with the `reduction` -clause

`reduction(operator:list)`

- Performs reduction on the (scalar) variables in list
- Private reduction variable is created for each thread's partial result
- Private reduction variable is initialized to operator's initial value
- After parallel region the reduction operation is applied to private variables and result is aggregated to the shared variable

OpenMP Work Sharing

Reductions within for loop:

```
int sum_variable = 0;

#pragma omp parallel for reduction(+: sum_variable)
for(int i=0; i<n; i++){
    sum_variable += ...;
    ...
}
```

Different reductions operators available like:

Operator Initial value

+	0
-	1
*	0

OpenMP Work Sharing

Linear algebra

Operation:

$$z_i = ax_i + y_i$$

- Copy ex03 as before:

```
[front01 ex02]$ cd ../  
[front01 tomp]$ cp -r /nvme/scratch/jfinkenrath/NCC_Training/ex03 .  
[front01 tomp]$ cd ex03
```

- Inspect, compile axpy.c, edit the submit script and run :

```
[front01 ex03]$ cc -std=c99 -fopenmp -o axpy axpy.c  
[front01 ex03]$ vi submit_ex03.sh  
...  
[front01 ex03]$ sbatch submit_ex03.sh  
[front01 ex03]$ more ex03.out  
t0 = 0.232100 sec, t1 = 0.232260 sec, diff z norm = 0.000000e+00
```

- Use an OpenMP pragma to parallelize the second occurrence of the main for loop

OpenMP Work Sharing

Linear algebra

Operation:

$$z_i = ax_i + y_i$$

- Use an OpenMP pragma to parallelize the second occurrence of the main for loop

It's also useful to report the total number of threads:

```
printf(" t0 = %lf sec, t1 = %lf sec, diff z)

#pragma omp parallel
{
    int nth = omp_get_num_threads();
    #pragma omp single
    printf(" nth = %2d, t0 = %lf sec, t1 = %lf)
}
```

OpenMP Work Sharing

Linear algebra

Operation:

$$z_i = ax_i + y_i$$

- Run for OMP_NUM_THREADS from 1,...,10. How does the runtime scale?

OpenMP Work Sharing

Linear algebra

Operation:

$$z_i = ax_i + y_i$$

- Run for OMP_NUM_THREADS from 1,...,10. How does the runtime scale?

```
[front01 ex03]$ vi submit_ex03.sh
...

for n in 1 2 3 4 5 6 7 8 9 10
do
OMP_NUM_THREADS=$n ./axpy $((32*1024*1024))
done
...
[front01 ex03]$ sbatch submit_ex03.sh
```

- How the speed up depend on the number of threads ?
- Does it change with the number of operation ?

OpenMP Work Sharing

Linear algebra

Dot product operation:

$$r = x^T y = \sum_{i=0}^{n-1} x_i y_i$$

- Copy ex04 as before:

```
[front01 ex03]$ cd ../  
[front01 tomp]$ cp -r /nvme/scratch/jfinkenrath/NCC_Training/ex04 .  
[front01 tomp]$ cd ex04
```

- Inspect, compile, and run xdoty.c :

```
[front01 ex04]$ cc -std=c99 -o xdoty xdoty.c  
[front01 ex04]$ vi submit_ex04.sh  
...  
./xdoty $((32*1024*1024))  
...  
[front01 ex04]$ sbatch submit_ex04.sh  
[front01 ex04]$ more ex04.out  
t0 = 0.172530 sec, t1 = 0.171624 sec, norms = 8.387960e+06, 8.387960e+06
```

- Use an OpenMP pragma to parallelize the second occurrence of the main for loop

OpenMP Work Sharing

Linear algebra

Dot product operation:

$$r = x^T y = \sum_{i=0}^{n-1} x_i y_i$$

- Now run for OMP_NUM_THREADS from 1,...,10.
- edit submit_ex04.sh via

```
for n in 1 2 3 4 5 6 7 8 9 10; do
    OMP_NUM_THREADS=$n ./xdoty $((32*1024*1024))
done
```

Does the code scale ?